

Helix Proxy — VPN-Aware Dynamic Routing

Revision: 1 **Last modified:** 2026-06-30T00:00:00Z **Status:** DESIGN — NOT YET IMPLEMENTED. This document describes the *intended* control-plane for VPN-aware dynamic routing. No code, container, or test for the dynamic mode exists in this repository yet; every behaviour below is written as a design intention (“will” / “is designed to”), never as a working claim (Constitution §11.4.6 — no false “implemented” assertions). **Authority:** Inherits the Helix Constitution submodule (constitution/Constitution.md) per §11.4.35. **Design source:** docs/superpowers/specs/2026-06-30-vpn-aware-proxy-extension-design.md.

1. Plain-language overview (for everyone)

Today the Helix Proxy can send **all** of its traffic through **one** VPN tunnel, or through **no** VPN at all. That is an all-or-nothing choice made once, when the service starts.

The **dynamic routing** extension is designed to make that choice **per destination** and **live**. In everyday terms, it will let you say things like:

- “Send requests for *server A* through *VPN tunnel 1*, and requests for *server B* through *VPN tunnel 2*.” Different destinations, different tunnels, **at the same time** — multiple tunnels running concurrently.
- “If the tunnel a destination needs is **down**, don’t hang and don’t silently leak my traffic — give me an **instant, clean ‘service unavailable’ (HTTP 503)** so I know exactly what happened.” When the tunnel comes back, requests start working again on their own — **no restart, no reconfiguration**.

Why this matters: it lets one proxy reach many private servers that each live behind a *different* VPN, it fails **gracefully and visibly** instead of leaving you guessing, and it recovers by itself. The design keeps the existing caching proxy doing its job, so web pages still load fast from the local cache.

This is a planned capability. None of it is built yet — see the Status line above.

2. The four operating modes

The three existing modes are **preserved** (they will not be removed — Constitution §11.4.122). The dynamic mode is the **new, additive** fourth mode.

Mode	How traffic is routed	Tunnels	Per-target choice	Down-tunnel behaviour	Status
vpn (existing)	All proxy traffic shares one VPN container's network namespace	One	No — all or nothing	Dropped / hung (no clean signal)	Shipped
host-vpn (existing)	Proxy runs on the host network and uses the host's own VPN	One (host's)	No	Depends on host VPN	Shipped
no-vpn (existing)	Direct to the internet, bridge network, local caching only	None	N/A	N/A	Shipped
dynamic (NEW)	Each destination is routed to a chosen VPN profile; a control-plane publishes live tunnel state	Many, concurrent	Yes — per target, with ordered failover tiers	Instant graceful 503 , auto-recovers	DESIGN ONLY — not built

The existing `dperson/openvpn-client` VPN is designed to be retained inside the dynamic mode as a legacy profile (marked deprecated, **not** removed).

3. Architecture (for engineers)

3.1 The out-of-byte-path principle

The single most important design decision: **the control-plane never sits in the request byte path**. Squid (HTTP/HTTPS + cache) and Dante (SOCKS5) keep doing all the actual proxying and caching, exactly as they do today. The new Go control-plane is **not** a reverse proxy — it is three cooperating roles that sit *beside* the data plane:

1. **config-compiler** — reads PostgreSQL and renders the static parts of the Squid / Dante configuration (per-tunnel `cache_peer` entries, Dante routes, a PAC file). It only touches the proxies for **structural** changes (a new tunnel or peer appears) via `reconfigure / SIGHUP` — **never** for a tunnel merely going up or down.
2. **health-publisher** — one watcher per tunnel. It will determine tunnel health from **data-plane facts** (gluetun control API, a `wg show` byte-transfer delta, and a live egress probe), never from “it is configured”, and write the result into Redis.
3. **external_acl helper** — a tiny Go binary that Squid calls **per request** through Squid’s native `external_acl_type` hook. It reads Redis and answers `OK tag=<tunnel>` (route this request to that tunnel’s peer) or `ERR` (which Squid turns into a branded `deny_info 503`).

This preserves Squid’s mature caching **and** adds dynamic per-target VPN selection plus a graceful 503 with **zero reconfigure** on every up/down flip.

3.2 The gluetun-per-profile model

Each VPN profile is designed to be **one gluetun container = one network namespace = one Redis status key** (`vpn:status:<profile>`). This maps 1:1 onto the existing `network_mode: service:<vpn>` pattern the project already uses, so it extends rather than reinvents the current approach. Kernel WireGuard is the preferred tunnel type (throughput, CPU, audit surface), with OpenVPN retained for compatibility.

3.3 PostgreSQL vs Redis — two stores, two jobs


- **PostgreSQL — the source of truth**. Durable configuration: VPN profiles, target hosts, proxy rules, ordered failover tiers,

users, and an audit log. The control-plane *reads* it to build configuration; it changes slowly and only via the admin API. Full schema and ER diagram: docs/DATABASE.md.

- **Redis — the live state and event bus.** Fast, ephemeral, per-request: `vpn:status:<profile>` (current up/down + handshake age + byte counters + egress IP, with a TTL so a **stale key is treated as down — fail-closed**), `vpn:events` (pub/sub so the helper and UI react instantly), and `route:<target>` (the compiler-resolved tunnel + tier + breaker state).

3.4 Architecture diagram

Diagram source: docs/diagrams/
dynamic_routing_architecture.mermaid (kept valid so it can later be export-validated per §11.4.168).



4. Request lifecycle

The intended path of a single HTTP/HTTPS request in dynamic mode:


1. **Client → Squid.** The client points at Squid (:53128) exactly as today.
2. **Squid → external_acl helper.** Before choosing an upstream, Squid calls the Go helper through `external_acl_type`, passing the request's Host.
3. **Helper → Redis.** The helper reads `route:<target>` and `vpn:status:<profile>` to learn `{target → tunnel, up?}`.
4. **Routing decision:**
 - **Tunnel up → helper returns OK tag=tun_<profile>;** Squid sends the request out through that tunnel's `cache_peer` (egress via the matching gluetun container), serving from cache where possible.
 - **Tunnel down / Redis stale → helper returns ERR;** Squid returns a **branded deny_info 503 (ERR_TUNNEL_DOWN)**. Crucially, the Squid process is **unchanged** — this is a clean L7 signal, not a crash.
5. **Auto-recovery — zero reconfigure.** When the tunnel returns, the health-publisher flips `vpn:status:<profile>` back to up and publishes `vpn:events`. The very next request gets OK from the

helper and succeeds. **No reconfigure, no SIGHUP, no restart** — only the per-request Redis read changed its answer.

This same Redis-driven flip is what delivers **circuit-breaker tier-failover**: if a target's primary tunnel is unhealthy, the helper is designed to fail over to the next configured tier before falling back to the 503.

4.1 Request sequence diagram

Diagram source: docs/diagrams/dynamic_routing_sequence.mermaid.



5. Failure modes (designed to fail closed)

The design's guiding rule is **fail-closed, never leak**:

- **Tunnel down** → branded 503, Squid PID unchanged (intended to be proven, not assumed).
- **Redis down / stale key** → fail-closed 503 — never fall through to a direct (leaking) request.
- **Circuit-breaker open for a target** → fail over to the next healthy tier, else 1.

- **Control-plane crash** → Squid keeps serving its last good configuration (degrade, do not die).

These are **intended** behaviours; each will require captured data-plane evidence (tcpdump, wg transfer delta, egress-IP probes) before it can be claimed to work — see the spec’s anti-bluff test strategy.

6. Cross-references

- **Design spec (full rationale, research mapping, test strategy):** docs/superpowers/specs/2026-06-30-vpn-aware-proxy-extension-design.md
- **Implementation plan (phasing P0-P12):** docs/superpowers/plans/2026-06-30-vpn-aware-proxy-extension-plan.md
- **Data model (PostgreSQL schema + ER diagram):** docs/DATABASE.md
- **Existing network modes (the three preserved modes):** docs/NETWORK_MODES.md
- **Existing system architecture:** docs/ARCHITECTURE.md
- **Caching behaviour (preserved under dynamic):** docs/CACHE.md
- **VPN configuration:** docs/VPN.md
- **Diagram sources:** docs/diagrams/dynamic_routing_architecture.mermaid, docs/diagrams/dynamic_routing_sequence.mermaid, docs/diagrams/database_er.mermaid

Document status reminder: the VPN-aware dynamic-routing control-plane is a design, not a shipped capability. This document will be updated to describe observed, evidence-backed behaviour only as each phase of the companion plan is implemented and verified.